

[Report 027] Databricks - AI Engineer - FDE (Forward Deployed Engineer) - Amsterdam

Score: 2.5/5
Culture Fit: Conditional
URL: https://databricks.com/company/careers/open-positions/job?gh_jid=8560401002
PDF: ❌
Legitimacy: High Confidence
Verification: confirmed (WebFetch)
Date: 2026-05-21

A) ROLE SUMMARY

Field	Detail
Archetype	AI Forward Deployed Engineer
Domain	Enterprise data & AI platform (GenAI, LLMOps, Databricks Intelligence Platform)
Function	Deploy + consult + build - customer-facing technical delivery of GenAI solutions
Seniority	All levels (Junior to Senior; JD title says "ALL LEVELS")
Remote	Remote Netherlands (preferred Amsterdam; also London, Madrid, Paris)
Team size	Not specified; part of the AI FDE EMEA team, reports to Senior Manager AI FDE EMEA
TL;DR	Client-facing GenAI engineer who builds and deploys production AI applications for Databricks enterprise customers - heavy data science / ML engineering requirements, advisory + delivery combined.

B) MATCH WITH CV

Requirements vs CV mapping

JD Requirement	CV Match	Strength
Building GenAI applications (RAG, multi-agent, Text2SQL, fine-tuning)	Self-hosted AI stack: Flowise, OpenWebUI, n8n, Supabase on Linux infrastructure. VCH AI compliance tools.	Partial - runs and integrates AI tools, but has not built production RAG or multi-agent systems from scratch using ML libraries
Deploying production-grade GenAI applications with evals	Self-hosted stack is in active use by students and researchers. No formal eval frameworks.	Partial - operational deployment yes, but not enterprise-grade with MLOps tooling
Extensive hands-on data science experience (pandas, scikit-learn, PyTorch)	Data Analysis & Visualization cert (R, Python). Dashboarding & Analysis course at InHolland. Oracle 11g BI deployments at KONGSBERG.	Weak - data literacy confirmed, but no evidence of pandas/scikit-learn/PyTorch at production level
Production ML deployments on AWS/Azure/GCP	Docker, Kubernetes, GitHub CI/CD, Linux self-hosted. No public cloud ML platform experience.	Weak - strong DevOps on-prem, but cloud ML platforms (SageMaker, AzureML, Vertex) not in CV
Databricks Intelligence Platform / Apache Spark experience	Not in CV	Gap
Quantitative degree or equivalent experience	Bachelor Business IT & Management (HvA, grade 8.7). Technical Informatics (HvA).	Partial - business informatics, not quantitative/STEM. Equivalent clause may cover it

Technical communication to diverse audiences	Expert: workshop design, train-the-trainer, KONGSBERG client training (BP, Statoil), university lecturing	Strong - this is a genuine superpower
Willingness to travel quarterly for customer visits	Remote-first EU is preferred. Quarterly travel is acceptable given fractional model.	Acceptable
Trusted technical advisor to customers	KONGSBERG: full BI implementations for BP (Jordan, Oman), Statoil. Client-facing technical delivery.	Moderate - real enterprise client delivery, but in 2012; different tech era

Gaps

Gap	Blocker?	Mitigation
No production ML / data science engineering (pandas, scikit-learn, PyTorch)	Hard blocker - this is listed as a core requirement for an "engineer" role	Would need 6-12 months of dedicated ML engineering work to close
No cloud ML platform experience (AWS SageMaker, Azure ML, GCP Vertex)	Hard blocker for enterprise AI deployment at Databricks customer level	Adjacent DevOps skills help but don't substitute
No Databricks/Spark/Delta Lake experience	Soft blocker (preferred, not required) but very expected given the company	Databricks has learning resources; 2-4 months to get functional
No formal eval frameworks or LLMOps tooling	Hard blocker - evals are explicitly required ("deploying production-grade GenAI applications with evaluation capabilities")	Ragas, DeepEval, PromptFlow are learnable; no production examples to show
No RAG/multi-agent system architecture from scratch	Hard blocker for credibility with enterprise customers at Databricks level	n8n workflows and Flowise are relevant but not at the same engineering depth

Summary: Christiaan is a practitioner and integrator of AI tools - he runs complex AI stacks for real users. But this role requires someone who builds the underlying ML systems, not someone who deploys and connects them. The "AI Forward Deployed" archetype in `_profile.md` is listed as "Primary" for Christiaan - but that assumes the FDE role is more advisory/implementation than data science engineering. This Databricks FDE leans heavily into ML engineering depth (PyTorch, scikit-learn, production deployments on cloud platforms). That's the critical gap.

C) LEVEL AND STRATEGY

Level detected: All levels explicitly stated; likely looking for mid-to-senior based on requirements (production ML, eval frameworks, enterprise client management)

Christiaan's natural level for this archetype: He maps to the advisory/facilitation half of this role well, but not the ML engineering half. In Databricks' framing, he would not pass technical screening for the engineering components.

"Sell senior without lying" plan:

- Lead with the KONGSBERG enterprise delivery track record (BP, Statoil - real enterprise client-facing technical delivery)
- Lead with AI stack operational depth (self-hosted, production, real users)
- Lead with technical communication and workshop/training capabilities as the differentiator
- But be honest: the data science engineering depth (ML libraries, cloud ML platforms) is not there

"If they downlevel me" - not applicable: The gap is not about level, it's about the wrong skill set for an engineering-forward FDE role.

Honest assessment: Christiaan should explore "Technical AI PM" or "Solutions Consultant" adjacent roles at Databricks, not the engineering FDE track. The advisory and communication skills are there; the ML engineering is not.

D) COMP AND DEMAND

Metric	Estimate	Source
Market range (AI FDE, Amsterdam, all levels)	€70,000-€130,000 base (junior to senior) + RSUs	Glassdoor Databricks FDE/Professional Services EMEA; Levels.fyi
Databricks compensation reputation	Strong base + RSU package; total comp can be 20-40% above base	Known industry data
Christiaan's target range	€65,000-€90,000	_profile.md
Fit	The lower end of the range overlaps with Christiaan's target if applying at junior/mid level	

Note: Compensation overlap exists at the lower end. But the technical requirements mean this is unlikely to progress to offer stage.

E) CUSTOMIZATION PLAN

Not recommended. Provided if user wishes to override:

#	Section	Current	Proposed change	Why
1	Summary	Lab builder / educator	Reframe as "applied AI practitioner building AI-powered workflows and research infrastructure"	Closest honest frame
2	Windesheim AI Stack	"Self-hosted AI sandbox"	Describe architecture explicitly: Docker-deployed, n8n/Flowise/OpenWebUI/Supabase, managing inference endpoints for student use	Shows production ops depth
3	KONGSBERG experience	Listed as BP/Statoil deployments	Lead with "enterprise client-facing technical delivery" - most relevant analog to FDE model	Real signal
4	VCH	Supply chain AI tools	Frame as "built AI compliance tooling for EUDR/CSRD - practical GenAI applications with real business users"	Honest applied AI proof
5	Data Analysis cert	Listed under certifications	Surface alongside Python skills to indicate ML literacy floor	Minimal but relevant

F) INTERVIEW PLAN

Abbreviated - role not recommended.

#	JD Requirement	STAR+R Story	Reflection
1	Client-facing technical delivery	KONGSBERG BP Jordan: full BI implementation - hardware, server config, software, dashboards, on-site training	Enterprise client delivery is real; tech era is old but pattern is identical
2	Building AI applications for real users	Windesheim AI stack: self-funded, Docker-deployed, production use by students and researchers	Shows initiative and production orientation; gap is depth of ML stack
3	Technical communication to diverse audiences	Blockchain Minor at InHolland: designed "Build Your Own Crypto" course, still running independently	Expert communicator of complex technical concepts
4	Travel and customer engagement	KONGSBERG: multiple countries (Jordan, Oman, Norway, Azerbaijan) in field service	Comfortable with travel for customer work

Red-flag questions:

- "Walk me through a RAG system you built end-to-end." → Honest answer: "I've deployed Flowise (which abstracts RAG), but I haven't built a custom RAG pipeline from scratch with LangChain or LlamaIndex. That's a gap I'd close quickly given the learning budget."

- "What's your experience with PyTorch or scikit-learn in production?" → Honest answer: "Data analysis in Python for research and coursework, but not production ML model training. My strength is at the application and integration layer."

G) POSTING LEGITIMACY

Assessment: High Confidence

Signal	Finding	Weight
Posting freshness	Active on Greenhouse API, requisition ID present	Positive
Apply button	Active	Positive
JD specificity	Highly specific - named technologies (RAG, multi-agent, Text2SQL, pandas, scikit-learn, PyTorch, AWS/Azure/GCP), named team structure (Senior Manager AI FDE EMEA)	Positive
Requirements realism	Consistent and realistic for a professional services engineering role at this level	Positive
Company hiring signals	Databricks consistently growing professional services capacity in EMEA	Positive
Role-company fit	Core business model - PS teams deliver against the platform	Positive

H) CULTURE, PERSONALITY & LONG-TERM FIT

H1 - Org Culture Assessment

Stage and type: Databricks is pre-IPO, high-growth. The AI FDE (Forward Deployed Engineer) team is a technical customer-facing PS function - part ML engineer, part consultant, part thought leader. The culture of this specific team tends to be high-energy, technically driven, and fast-moving. FDE roles at companies like Databricks attract people who want to be at the intersection of cutting-edge AI research and enterprise delivery.

Culture signals: The JD mentions "present at conferences and serve as thought leader," "fueling curiosity for the latest trends in GenAI," "collaborate with product and engineering on roadmap priorities." This is a technically ambitious, community-visible, curiosity-driven culture within the commercial PS wrapper.

Physical/work environment: Remote Netherlands (preferred Amsterdam). Quarterly customer travel. The base environment is remote, with periodic client sites and Databricks events.

Hackerspace vision match: Moderate. The technical curiosity culture, thought leadership mandate, and frontier AI focus (GenAI, LLMOps, multi-agent) align with the hackerspace energy. The "new territory" element is genuinely present - this is a fast-moving frontier. But the enterprise client delivery context (Fortune 500 customers, production ML) requires sustained technical depth that is currently missing. Score: 6/10 against the 2-year vision if the skill gap were resolved.

H2 - Personality Fit Table

Dimension	What the role requires	Christiaan's profile	Match
Starter vs. finisher	Each new customer engagement has a 0-to-1 element (new problem, new architecture)	ENFP-T starter	Good match - the per-engagement cycle rewards the starter energy
Variety vs. repetition	High variety: new customers, new use cases, new GenAI techniques	Thrives on variety	Strong match - FDE is one of the highest-variety PS roles
Autonomy level	High autonomy within each engagement	Needs ownership	Good match
Build vs. maintain	Building GenAI applications for each customer - build mandate per engagement	Builder identity	Strong match - each deployment is a build
Fractional fit	Full-time implied; but quarterly travel model is lighter than 50%	Optimal: 2-3 concurrent 0.4-0.6 FTE	Structural conflict but lighter than Elastic/Munich

Energy spike tolerance	Technical curiosity culture; conference thought leadership rewards visible energy	Front-loaded high energy	Good fit - the culture rewards the energy spike
Team size	EMEA FDE team - likely a small, niche technical team	Small team, high trust	Good fit

Flags: fractional conflict; ML engineering skill gap creates critical imposter risk in this specific team culture.

H3 - Values & Mission Alignment

The AI FDE culture - building at the frontier, sharing knowledge publicly, solving real enterprise problems with cutting-edge techniques - is genuinely aligned with Christiaan's identity as a builder and educator. The thought leadership and conference presence component is an intrinsic motivator (guest lectures unpaid, knowledge sharing is part of his identity). The "fueling curiosity" framing is exactly the kind of culture he would thrive in.

The problem is the ML engineering gap creates a values mismatch of a different kind: being in a room where colleagues are discussing PyTorch fine-tuning and custom RAG pipelines from scratch, without the depth to contribute at that level, would generate chronic imposter syndrome and undermine the energy and authenticity that are Christiaan's strengths.

Verdict: Mission alignment is high. Skills alignment is low. The culture would be right; the current technical level does not match it.

H4 - Energy Analysis

Energizers:

- Per-engagement variety: each new customer is a new problem
- Frontier AI domain: GenAI, LLMOps, multi-agent - genuinely exciting territory
- Thought leadership and conference presence: public knowledge sharing is intrinsic
- Small EMEA team with high technical curiosity culture
- Build mandate per engagement: every FDE project has something to create
- Customer communication component: explaining complex AI to diverse audiences is a genuine strength

Drains:

- ML engineering depth gap: the gap between "runs AI tools" and "builds ML systems from scratch" creates sustained anxiety in a team of ML engineers
- Production ML accountability: being responsible for production GenAI systems you did not build from scratch creates a specific kind of stress
- Full-time exclusivity: blocks the multi-engagement model
- The ramp time to close the ML engineering gap is 6-12 months minimum - that is a long period of imposter experience in a technically scrutinizing team

H5 - 18-Month Durability Test

Rating: Short-term fit IF the skill gap is closed first. Not viable at current skill level.

If Christiaan spent 6-12 months deliberately building ML engineering depth (LangChain, PyTorch basics, RAG from scratch, LLMOps tooling), this role would become a strong fit: the culture is right, the variety is right, the domain is right, the thought leadership mandate is right. But applying now and ramping on the skills on the job would be a very difficult experience in a technically exacting peer environment.

If hired at a junior level with an explicit learning mandate, the 18-month trajectory could be positive - but that requires the company to accept a practitioner-level hire on the engineering track.

H6 - Personality-Specific Red Flags

Red flag	Present?	Severity	Context
Role is primarily maintenance/ops	No	N/A	Each engagement has a build mandate
Solo sales of an unready product	No	N/A	Not a sales role
Full-time exclusivity expected	Yes	Medium	Full-time implied; lighter travel model than some

Large corporate, slow decision cycles	No	N/A	Databricks moves fast; FDE team is agile
Ceremony-only role	No	N/A	Not applicable
Governance in decline	No	N/A	Not applicable

H7 - Fit Verdict

Case FOR: The FDE culture at Databricks is arguably the best cultural environment of all ten roles evaluated in this batch: frontier AI, small technical team, high variety, thought leadership, build mandate per engagement, customer communication. The personality fit on every dimension except the skill gap is genuinely strong. If Christiaan closes the ML engineering gap, this becomes a 4.0+ role.

Case AGAINST (watch for): The ML engineering gap is not a framing problem - it is a technical depth problem that a technically sophisticated peer group will surface quickly. Joining this team at current skill level would place Christiaan in a chronically uncomfortable position. The culture match being high makes the skill gap more painful, not less, because the culture is exactly what would motivate Christiaan to perform - and the performance ceiling is set by a skill he does not yet have.

Score cap note: The 2.5/5 score reflects the skill gaps. The culture fit assessment confirms the score is appropriate: strong culture potential, low current viability. The right recommendation is "build toward this archetype, not apply now."

Overall personality/culture fit: Conditional (strong culture match; conditional on ML engineering skill development)

KEYWORDS EXTRACTED

Forward Deployed Engineer, FDE, GenAI, RAG, multi-agent, Text2SQL, fine-tuning, LLMOps, pandas, scikit-learn, PyTorch, AWS, Azure, GCP, Databricks Intelligence Platform, Apache Spark, production ML, professional services, customer-facing, EMEA, evaluation frameworks, AI applications

JOB DESCRIPTION (ARCHIVED)

Archived 2026-05-21 - original posting may be removed after offer closes

Title: AI Engineer - FDE (Forward Deployed Engineer) (ALL LEVELS)

Requisition ID: CSQ327R197

Location: Netherlands - Remote (preferred: London, Madrid, Paris, Amsterdam)

Department: Professional Services Operations

Recruiter: Dina Hussain

Reports to: Senior Manager - AI FDE, EMEA

Mission:

The AI FDE team delivers professional services to help customers build and productionize AI applications. The role involves working with customers, teammates, and fueling curiosity for the latest trends in GenAI, LLMOps, and ML.

Key Responsibilities:

- Develop GenAI solutions using latest research techniques
- Own production rollouts of consumer and internally facing GenAI applications
- Serve as trusted technical advisor to customers
- Present at conferences and serve as thought leader
- Collaborate with product and engineering teams on roadmap priorities

Requirements:

- Experience building GenAI applications (RAG, multi-agent systems, Text2SQL, fine-tuning)
- Expertise deploying production-grade GenAI applications with evaluation capabilities
- Extensive hands-on data science experience with pandas, scikit-learn, PyTorch
- Production ML deployments on AWS, Azure, or GCP
- Quantitative degree or equivalent experience
- Technical communication skills

- Willingness to travel quarterly for customer visits

Preferred:

- Databricks Intelligence Platform and Apache Spark experience

About Databricks: Over 10,000 organizations including Fortune 500 companies rely on the Databricks Data Intelligence Platform for unified data, analytics, and AI capabilities.